

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

[1]. Reading Data

[1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

In [1]:

```
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
```

```
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os
```

In [2]:

```
# using SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000""", con)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 """, con)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(0).
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)
```

Number of data points in our data (525814, 10)

Out[2]:

id	ProductId	Userid	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	1	1303862400	Good Quality Dog Food
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	0	1346976000	Not as Advertised
			Natalia Corres						"Delight"

2	3	B000LQOCH0	ABXLMWJIXXAIN	Corres	HelpfulnessNumerator	HelpfulnessDenominator	Score	1219017600	Design
Id	ProductId		UserId	ProfileName				Time	Summary

In [3]:

```
display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

In [4]:

```
print(display.shape)
display.head()
```

(80668, 7)

Out[4]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...	2
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

In [5]:

```
display[display['UserId']=='AZY10LLTJ71NX']
```

Out[5]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommended to try green tea extract to ...	5

In [6]:

```
display['COUNT(*)'].sum()
```

Out[6]:

393063

[2] Exploratory Data Analysis

[2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

In [7]:

```
display= pd.read_sql_query("""
SELECT *
```

```
SELECT
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

Out[7]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summ
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADRA VANII WAFE
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADRA VANII WAFE
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADRA VANII WAFE
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADRA VANII WAFE
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADRA VANII WAFE

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

In [8]:

```
#Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
```

In [9]:

```
#Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first', inplace=False)
final.shape
```

Out[9]:

(364173, 10)

In [10]:

```
#Checking to see how much % of data still remains
(final.shape[0]-1)/(filtered_data.shape[0]-1)*100
```

```
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

Out[10]:

69.25890143662969

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

In [11]:

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

Out[11]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	1	5	1224892800	Bought This for My Son at College
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	2	4	1212883200	Pure cocoa taste with crunchy almonds inside

In [12]:

```
final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

In [13]:

```
#Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

(364171, 10)

Out[13]:

```
1    307061
0     57110
Name: Score, dtype: int64
```

[3] Preprocessing

[3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags

2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

In [14]:

```
# printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along and he always can sing the refrain. he's learned about whales, India, drooping roses: i love all the new words this book introduces and the silliness of it all. this is a classic book i am willing to bet my son will STILL be able to recite from memory when he is in college

I was really looking forward to these pods based on the reviews. Starbucks is good, but I prefer bolder taste.... imagine my surprise when I ordered 2 boxes - both were expired! One expired back in 2005 for gosh sakes. I admit that Amazon agreed to credit me for cost plus part of shipping, but geez, 2 years expired!!! I'm hoping to find local San Diego area shoppe that carries pods so that I can try something different than starbucks.

Great ingredients although, chicken should have been 1st rather than chicken broth, the only thing I do not think belongs in it is Canola oil. Canola or rapeseed is not something a dog would ever find in nature and if it did find rapeseed in nature and eat it, it would poison them. Today's Food industries have convinced the masses that Canola oil is a safe and even better oil than olive or virgin coconut, facts though say otherwise. Until the late 70's it was poisonous until they figured out a way to fix that. I still like it but it could be better.

Can't do sugar. Have tried scores of SF Syrups. NONE of them can touch the excellence of this product.

Thick, delicious. Perfect. 3 ingredients: Water, Maltitol, Natural Maple Flavor. PERIOD. No chemicals. No garbage.

Have numerous friends & family members hooked on this stuff. My husband & son, who do NOT like "sugar free" prefer this over major label regular syrup.

I use this as my SWEETENER in baking: cheesecakes, white brownies, muffins, pumpkin pies, etc... Unbelievably delicious...

Can you tell I like it? :)

In [15]:

```
# remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along and he always can sing the refrain. he's learned about whales, India, drooping roses: i love all the new words this book introduces and the silliness of it all. this is a classic book i am willing to bet my son will STILL be able to recite from memory when he is in college

In [16]:

```
# https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-tags-from-an-element
```

```

from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)

```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along and he always can sing the refrain. he's learned about whales, India, drooping roses: i love all the new words this book introduces and the silliness of it all. this is a classic book i am willing to bet my son will STILL be able to recite from memory when he is in college

I was really looking forward to these pods based on the reviews. Starbucks is good, but I prefer bolder taste.... imagine my surprise when I ordered 2 boxes - both were expired! One expired back in 2005 for gosh sakes. I admit that Amazon agreed to credit me for cost plus part of shipping, but geez, 2 years expired!!! I'm hoping to find local San Diego area shoppe that carries pods so that I can try something different than starbucks.

Great ingredients although, chicken should have been 1st rather than chicken broth, the only thing I do not think belongs in it is Canola oil. Canola or rapeseed is not something a dog would ever find in nature and if it did find rapeseed in nature and eat it, it would poison them. Today's Food industries have convinced the masses that Canola oil is a safe and even better oil than olive or virgin coconut, facts though say otherwise. Until the late 70's it was poisonous until they figured out a way to fix that. I still like it but it could be better.

Can't do sugar. Have tried scores of SF Syrups. NONE of them can touch the excellence of this product. Thick, delicious. Perfect. 3 ingredients: Water, Maltitol, Natural Maple Flavor. PERIOD. No chemicals. No garbage. Have numerous friends & family members hooked on this stuff. My husband & son, who do NOT like "sugar free" prefer this over major label regular syrup. I use this as my SWEETENER in baking: cheesecakes, white brownies, muffins, pumpkin pies, etc... Unbelievably delicious...Can you tell I like it? :)

In [17]:

```

# https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase

```

In [18]:

```

sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)

```

Great ingredients although, chicken should have been 1st rather than chicken broth, the only thing

Great ingredients although, chicken should have been 1st rather than chicken broth, the only thing I do not think belongs in it is Canola oil. Canola or rapeseed is not something a dog would ever find in nature and if it did find rapeseed in nature and eat it, it would poison them. Today is Food industries have convinced the masses that Canola oil is a safe and even better oil than olive or virgin coconut, facts though say otherwise. Until the late 70 is it was poisonous until they figured out a way to fix that. I still like it but it could be better.

=====

In [19]:

```
#remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along and he always can sing the refrain. he's learned about whales, India, drooping roses: i love all the new words this book introduces and the silliness of it all. this is a classic book i am willing to bet my son will STILL be able to recite from memory when he is in college

In [20]:

```
#remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

Great ingredients although chicken should have been 1st rather than chicken broth the only thing I do not think belongs in it is Canola oil Canola or rapeseed is not something a dog would ever find in nature and if it did find rapeseed in nature and eat it it would poison them Today is Food industries have convinced the masses that Canola oil is a safe and even better oil than olive or virgin coconut facts though say otherwise Until the late 70 is it was poisonous until they figured out a way to fix that I still like it but it could be better

In [21]:

```
# https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have been removed in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", \
               "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', \
               'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them', 'their', \
               'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'those', \
               'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', 'do', 'does', \
               'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'while', 'of', \
               'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during', 'before', 'after', \
               'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under', 'again', 'further', \
               'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'each', 'few', 'more', \
               'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very', 's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll', 'm', 'o', 're', \
               've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn', \
               "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn't", 'mustn', \
               "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', "wasn't", 'weren', "weren't", \
               'won', "won't", 'wouldn', "wouldn't"])
```

In [22]:

```
# Combining all the above students
```



```

from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentence.strip())

```

100%|██████████| 364171/364171 [02:38<00:00, 2304.27it/s]

In [23]:

```
preprocessed_reviews[1500]
```

Out[23]:

'great ingredients although chicken rather chicken broth thing not think belongs canola oil canola rapeseed not someting dog would ever find nature find rapeseed nature eat would poison today food industries convinced masses canola oil safe even better oil olive virgin coconut facts though say otherwise late poisonous figured way fix still like could better'

[3.2] Preprocessing Review Summary

In [24]:

```

## Summary preprocessing
from tqdm import tqdm
preprocessed_summary = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Summary'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_summary.append(sentence.strip())

```

6%|███████| 22243/364171 [00:03<00:55, 6113.33it/s]/home/lab12/anaconda3/lib/python3.7/site-packages/bs4/__init__.py:273: UserWarning: "b'.'" looks like a filename, not markup. You should probably open this file and pass the filehandle into BeautifulSoup.
' Beautiful Soup.' % markup)
27%|███████| 97851/364171 [00:16<00:43, 6121.19it/s]/home/lab12/anaconda3/lib/python3.7/site-packages/bs4/__init__.py:273: UserWarning: "b'.'" looks like a filename, not markup. You should probably open this file and pass the filehandle into BeautifulSoup.
' Beautiful Soup.' % markup)
27%|███████| 98464/364171 [00:16<00:43, 6077.36it/s]/home/lab12/anaconda3/lib/python3.7/site-packages/bs4/__init__.py:273: UserWarning: "b'.'" looks like a filename, not markup. You should probably open this file and pass the filehandle into BeautifulSoup.
' Beautiful Soup.' % markup)
59%|███████| 216616/364171 [00:36<00:24, 5989.05it/s]/home/lab12/anaconda3/lib/python3.7/site-packages/bs4/__init__.py:273: UserWarning: "b'.'" looks like a filename, not markup. You should probably open this file and pass the filehandle into BeautifulSoup.
' Beautiful Soup.' % markup)
97%|███████| 354825/364171 [00:59<00:01, 5985.96it/s]/home/lab12/anaconda3/lib/python3.7/site-packages/bs4/__init__.py:273: UserWarning: "b'.'" looks like a filename, not markup. You should probably open this file and pass the filehandle into BeautifulSoup.
' Beautiful Soup.' % markup)
100%|██████████| 364171/364171 [01:00<00:00, 5993.20it/s]

In [23]:

```
final['Cleaned_text'] = preprocessed_reviews
```

In [24]:

```
### Sort data according to time series
final.sort_values('Time', inplace=True)
```

In [25]:

```
### Taking 100k samples
final_100k = final.sample(n=100000)
```

In [26]:

```
x = final_100k['Cleaned_text']
x.size
```

Out[26]:

100000

In [27]:

```
y = final_100k['Score']
y.size
```

Out[27]:

100000

In [28]:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=42)
```

[4] Featurization

[4.1] BAG OF WORDS

In [29]:

```
#BoW
count_vect = CountVectorizer() #in scikit-learn
x_train_bow = count_vect.fit_transform(x_train)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

x_test_bow = count_vect.transform(x_test)
print("the type of count vectorizer ", type(x_test_bow))
print("the shape of out text BOW vectorizer ", x_test_bow.get_shape())
print("the number of unique words ", x_test_bow.get_shape()[1])
```

some feature names ['aa', 'aaa', 'aaaa', 'aaaaa', 'aaaaaaaaaa', 'aaaaaaahhhhhh', 'aaaaaaahhh', 'aaaaaaahhhh', 'aaaaaaand', 'aaaaaaawwwwwwww']

=====

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>

the shape of out text BOW vectorizer (30000, 50937)

the number of unique words 50937

In [30]:

```
from sklearn import preprocessing
x_train_bow = preprocessing.normalize(x_train_bow)
x_test_bow = preprocessing.normalize(x_test_bow)
```

[4.2] Bi-Grams and n-Grams.

In [32]:

```
#bi-gram, tri-gram and n-gram

#removing stop words like "not" should be avoided before building n-grams
# count_vect = CountVectorizer(ngram_range=(1,2))
# please do read the CountVectorizer documentation http://scikit-learn.org/stable/modules/generated/sklearn.feature\_extraction.text.CountVectorizer.html

# you can choose these numebrs min_df=10, max_features=5000, of your choice
count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
final_bigram_counts = count_vect.fit_transform(x_train)
print("the type of count vectorizer ",type(final_bigram_counts))
print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_bigram_counts.get_shape()[1])
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (70000, 5000)
the number of unique words including both unigrams and bigrams 5000
```

[4.3] TF-IDF

In [31]:

```
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
x_train_tfidf = tf_idf_vect.fit_transform(x_train)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names()[0:10])
print('='*50)

x_test_tfidf = tf_idf_vect.transform(x_test)
print("the type of count vectorizer ",type(x_test_tfidf))
print("the shape of out text TFIDF vectorizer ",x_test_tfidf.get_shape())
print("the number of unique words including both unigrams and bigrams ", x_test_tfidf.get_shape()[1])
```

```
some sample features(unique words in the corpus) ['abandon', 'abandoned', 'abdominal', 'ability',
'able', 'able buy', 'able chew', 'able cut', 'able drink', 'able eat']
=====
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (30000, 40848)
the number of unique words including both unigrams and bigrams 40848
```

In [32]:

```
from sklearn import preprocessing
x_train_tfidf = preprocessing.normalize(x_train_tfidf)
x_test_tfidf = preprocessing.normalize(x_test_tfidf)
```

[4.4] Word2Vec

In [33]:

```
# Train your own Word2Vec model using your own text corpus
i=0
list_of_sentence_train=[]
for sentence in x_train:
    list_of_sentence_train.append(sentence.split())
```

In [34]:

```
# Test your own Word2Vec model using your own text corpus
i=0
list_of_sentence_test=[]
for sentence in x_test:
    list_of_sentence_test.append(sentence.split())
```

In [35]:

```
# Using Google News Word2Vectors

# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >12G of ram
# we will provide a pickle file wich contains a dict ,
# and it contains all our courpus words as keys and model[word] as values
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.

# http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.Wl7SRFAzZPY
# you can comment this whole cell
# or change these variable according to your need

is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occurred atleast 5 times
    w2v_model=Word2Vec(list_of_sentence_train,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin', binary=True)
    else:
        print("you don't have google's word2vec file, keep want_to_train_w2v = True, to train your own w2v ")
```

```
[('awesome', 0.8440885543823242), ('terrific', 0.825779378414154), ('wonderful',
0.8151693344116211), ('fantastic', 0.8132066130638123), ('excellent', 0.8085969686508179),
('good', 0.8001235127449036), ('perfect', 0.7367382645606995), ('amazing', 0.7261995077133179), ('
fabulous', 0.7111331224441528), ('incredible', 0.6996779441833496)]
=====
[('nastiest', 0.7835044264793396), ('tastiest', 0.7758582830429077), ('best', 0.7460561990737915),
('greatest', 0.7407208681106567), ('saltiest', 0.6663205027580261), ('superior',
0.6384615302085876), ('disgusting', 0.6267989277839661), ('healthiest', 0.5846405029296875),
('grossest', 0.5758904218673706), ('smoothest', 0.5751399397850037)]
```

In [36]:

```
w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])
```

```
number of words that occurred minimum 5 times 16192
sample words ['pure', 'toxic', 'garbage', 'eat', 'poison', 'want', 'metals', 'go', 'ahead',
'buy', 'junk', 'feed', 'look', 'high', 'ratings', 'apparently', 'people', 'love', 'eating', 'not',
'harmful', 'would', 'sodium', 'mercury', 'phosphate', 'enjoy', 'got', 'oil', 'free', 'ordered', 's
tuff', 'started', 'okay', 'based', 'smell', 'think', 'turned', 'rancid', 'time', 'let', 'tell', 'c
ooking', 'coconut', 'no', 'picnic', 'edit', 'year', 'later', 'may', 'lowered']
```

[4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

In [37]:

```
# average Word2Vec
# compute average word2vec for each review.
```

```

sent_vectors_train = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sentence_train): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to change this
    # to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_train.append(sent_vec)
print(len(sent_vectors_train))
print(len(sent_vectors_train[0]))

```

100%|██████████| 70000/70000 [03:16<00:00, 356.77it/s]

70000

50

In [38]:

```

# average Word2Vec
# compute average word2vec for each review.
sent_vectors_test = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sentence_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to change this
    # to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_test.append(sent_vec)
print(len(sent_vectors_test))
print(len(sent_vectors_test[0]))

```

100%|██████████| 30000/30000 [01:19<00:00, 377.26it/s]

30000

50

In [39]:

```

x_train_avgw2v=sent_vectors_train
x_test_avgw2v=sent_vectors_test

```

In [40]:

```

from sklearn import preprocessing
x_train_avgw2v = preprocessing.normalize(x_train_avgw2v)
x_test_avgw2v = preprocessing.normalize(x_test_avgw2v)

```

[4.4.1.2] TFIDF weighted W2v

In [41]:

```

# S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer(min_df=10, max_features=500)
tf_idf_matrix = model.fit_transform(x_train)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

```

In [42]:

```
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors_train = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sentence_train): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors_train.append(sent_vec)
    row += 1
```

100%|██████████| 70000/70000 [03:25<00:00, 341.30it/s]

In [43]:

```
tfidf_sent_vectors_test = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sentence_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors_test.append(sent_vec)
    row += 1
```

100%|██████████| 30000/30000 [01:29<00:00, 334.85it/s]

In [44]:

```
x_train_tfidf2v = tfidf_sent_vectors_train
x_test_tfidf2v = tfidf_sent_vectors_test
```

In [45]:

```
from sklearn import preprocessing
x_train_tfidf2v = preprocessing.normalize(x_train_tfidf2v)
x_test_tfidf2v = preprocessing.normalize(x_test_tfidf2v)
```

[5] Assignment 9: Random Forests

1. Apply Random Forests & GBDT on these feature sets

GET IT DONE: Random Forests and Gradient Boosting Decision Trees (GBDT)

- **SET 1**: Review text, preprocessed one converted into vectors using (BOW)
- **SET 2**: Review text, preprocessed one converted into vectors using (TFIDF)
- **SET 3**: Review text, preprocessed one converted into vectors using (AVG W2v)
- **SET 4**: Review text, preprocessed one converted into vectors using (TFIDF W2v)

2. The hyper parameter tuning (Consider two hyperparameters: `n_estimators` & `max_depth`)

- Find the best hyper parameter which will give the maximum [AUC](#) value
- Find the best hyper parameter using k-fold cross validation or simple cross validation data
- Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task of hyperparameter tuning

3. Feature importance

- Get top 20 important features and represent them in a word cloud. Do this for BOW & TFIDF.

4. Feature engineering

- To increase the performance of your model, you can also experiment with with feature engineering like :
 - Taking length of reviews as another feature.
 - Considering some features from review summary as well.

5. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure

with X-axis as `n_estimators`, Y-axis as `max_depth`, and Z-axis as **AUC Score** , we have given the notebook which explains how to plot this 3d plot, you can find it in the same drive [3d_scatter_plot.ipynb](#)

(or)

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure

[seaborn heat maps](#) with rows as `n_estimators`, columns as `max_depth`, and values inside the cell representing **AUC Score**

- You choose either of the plotting techniques out of 3d plot or heat map
- Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.
- Along with plotting ROC curve, you need to print the [confusion matrix](#) with predicted and original labels of test data points. Please visualize your confusion matrices using [seaborn heatmaps](#).

6. Conclusion

- [You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this prettytable library link](#)

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method `fit_transform()` on you train data, and apply the method `transform()` on cv/test data.
4. For more details please go through this [link](#).

[5.1] Applying RF

[5.1.1] Applying Random Forests on BOW, **SET 1**

In [46]:

```
## find hyperparameter using roc_auc score and plot AUC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]
```

```


#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for d in max_depth_values:
    clf = RandomForestClassifier(max_depth=d)
    clf.fit(x_train_bow, y_train)
    y_train_pred = clf.predict_proba(x_train_bow)[:,-1]
    y_test_pred = clf.predict_proba(x_test_bow)[:,-1]

    training_scores.append(roc_auc_score(y_train,y_train_pred))
    cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(max_depth_values, training_scores, label='Train AUC')
plt.plot(max_depth_values, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("max_depth: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```



In [45]:

```

## find hyperparameter using roc_auc score and plot AUC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
base_learners = [40,60,80,100,120]

#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for n in base_learners:
    clf = RandomForestClassifier(n_estimators=n)
    clf.fit(x_train_bow, y_train)
    y_train_pred = clf.predict_proba(x_train_bow)[:,-1]
    y_test_pred = clf.predict_proba(x_test_bow)[:,-1]


    training_scores.append(roc_auc_score(y_train,y_train_pred))
    cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(base_learners, training_scores, label='Train AUC')
plt.plot(base_learners, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("n_estimator: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```

ERROR PLOTS





In [47]:

```
### Using GridsearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV

base_learners = [40,60,80,100,120]
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

param_grid = {'n_estimators': base_learners , "max_depth":max_depth_values}
RFC = RandomForestClassifier(max_features='sqrt')
model = GridSearchCV(RFC, param_grid, scoring = 'roc_auc', cv=3 , n_jobs = -1,pre_dispatch=2)
model.fit(x_train_bow, y_train)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(x_test_bow, y_test))
```

```
Model with best parameters :
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=1000, max_features='sqrt', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=120, n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
Accuracy of the model : 0.9285171888288075
```

In [46]:

```
#### 3D Scatter Plot
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
import numpy as np

x1 = [40,60,80,100,120]
y1 = [1, 5, 10, 50, 100]
z1 = training_scores

x2 = [40,60,80,100,120]
y2 = [1, 5, 10, 50, 100]
z2 = cv_scores
```

In [47]:

```
# https://plot.ly/python/3d-axes/
trace1 = go.Scatter3d(x=x1,y=y1,z=z1, name = 'train')
trace2 = go.Scatter3d(x=x2,y=y2,z=z2, name = 'Cross validation')
data = [trace1, trace2]

layout = go.Layout(scene = dict(
    xaxis = dict(title='n_estimators'),
    yaxis = dict(title='max_depth'),
    zaxis = dict(title='AUC'),))

fig = go.Figure(data=data, layout=layout)
offline.iplot(fig, filename='3d-scatter-colorscale')
```

Observation: Using GridsearchCV max_depth is 1000 and n_estimators is 120.

ROC Curve using false positive rate versus true positive rate

In [48]:


```
# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
from sklearn.metrics import roc_curve, auc
from sklearn.ensemble import RandomForestClassifier
clf = RandomForestClassifier(max_depth=1000, n_estimators=120)
clf.fit(x_train_bow, y_train)
y_pred = clf.predict(x_test_bow)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs


train_fpr, train_tpr, thresholds = roc_curve(y_train, clf.predict_proba(x_train_bow)[:,1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, clf.predict_proba(x_test_bow)[:,1])

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.show()

print("="*100)

from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, clf.predict(x_train_bow)))
print("Test confusion matrix")
print(confusion_matrix(y_test, clf.predict(x_test_bow)))
```





Train confusion matrix

```
[[10786   26]
 [    0 59188]]
```

Test confusion matrix

```
[[  918  3917]
 [   46 25119]]
```

In [49]:


```
# Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm
```

Out[49]:

```
array([[  918,  3917],
       [   46, 25119]])
```

In [50]:

```
# plot confusion matrix to describe the performance of classifier.
import seaborn as sns
class_label = ["negative", "positive"]
df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
sns.heatmap(df_cm, annot = True, fmt = "d")
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```




[5.1.2] Wordcloud of top 20 important features from SET 1

In [51]:

```
# Calculate feature importances from decision trees
importances = clf.feature_importances_

# Sort feature importances in descending order
indices = np.argsort(importances)[::-1][:20]

# Rearrange feature names so they match the sorted feature importances
names = count_vect.get_feature_names()
```





[5.1.3] Applying Random Forests on TFIDF, SET 2

In [48]:


```
## find hyperparameter using roc_auc score and plot AUC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for d in max_depth_values:
    clf = RandomForestClassifier(max_depth=d)
    clf.fit(x_train_tfidf, y_train)
    y_train_pred = clf.predict_proba(x_train_tfidf)[:,1]
    y_test_pred = clf.predict_proba(x_test_tfidf)[:,1]

    training_scores.append(roc_auc_score(y_train,y_train_pred))
    cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(max_depth_values, training_scores, label='Train AUC')
plt.plot(max_depth_values, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.show()
```



In [55]:

```
## find hyperparameter using roc_auc score and plot AUC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
base_learners = [40,60,80,100,120]

#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for n in base_learners:
    clf = RandomForestClassifier(n_estimators=n)
    clf.fit(x_train_tfidf, y_train)
    y_train_pred = clf.predict_proba(x_train_tfidf)[:,1]
```


```

y_test_pred = clf.predict_proba(x_test_tfidf)[:,-1]

training_scores.append(roc_auc_score(y_train,y_train_pred))
cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(base_learners, training_scores, label='Train AUC')
plt.plot(base_learners, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("n_estimator: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```



In [51]:

```

## Using GridsearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV

base_learners = [40,60,80,100,120]
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

param_grid = {'n_estimators': base_learners , "max_depth":max_depth_values}
RFC = RandomForestClassifier(max_features='sqrt')
model = GridSearchCV(RFC, param_grid, scoring = 'accuracy', cv=3 , n_jobs = -1,pre_dispatch=2)
model.fit(x_train_tfidf, y_train)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(x_test_tfidf, y_test))

```

Model with best parameters :

```

RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=500, max_features='sqrt', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=40, n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)

```

Accuracy of the model : 0.8906666666666667

In [56]:

```

#### 3D Scatter Plot
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
import numpy as np

x1 = [40,60,80,100,120]
y1 = [1, 5, 10, 50, 100]
z1 = training_scores

x2 = [40,60,80,100,120]
y2 = [1, 5, 10, 50, 100]
z2 = cv_scores

```

In [57]:

```
# https://plot.ly/python/3d-axes/
trace1 = go.Scatter3d(x=x1,y=y1,z=z1, name = 'train')
trace2 = go.Scatter3d(x=x2,y=y2,z=z2, name = 'Cross validation')
data = [trace1, trace2]

layout = go.Layout(scene = dict(
    xaxis = dict(title='n_estimators'),
    yaxis = dict(title='max_depth'),
    zaxis = dict(title='AUC'),))

fig = go.Figure(data=data, layout=layout)
offline.iplot(fig, filename='3d-scatter-colorscale')
```

Observation: Using GridsearchCV max_depth is 500 and n_estimators is 40.

In []:

```
### ROC Curve using false positive rate versus true positive rate
```

In [52]:

```
# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
from sklearn.metrics import roc_curve, auc
clf = RandomForestClassifier(max_depth=500,n_estimators=40)
clf.fit(x_train_tfidf, y_train)
y_pred = clf.predict(x_test_tfidf)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs


train_fpr, train_tpr, thresholds = roc_curve(y_train, clf.predict_proba(x_train_tfidf)[:,-1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, clf.predict_proba(x_test_tfidf)[:,-1])

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
```

```
plt.show()

print("="*100)

from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, clf.predict(x_train_tfidf)))
print("Test confusion matrix")
print(confusion_matrix(y_test, clf.predict(x_test_tfidf)))
```



```
=====

Train confusion matrix
[[10962   55]
 [    0 58983]]
Test confusion matrix
[[ 1515   3137]
 [   131 25217]]
```

In [53]:


```
# Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm
```

Out[53]:

```
array([[ 1515,   3137],
       [    131, 25217]])
```

In [54]:

```
# plot confusion matrix to describe the performance of classifier.
import seaborn as sns
class_label = ["negative", "positive"]
df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
sns.heatmap(df_cm, annot = True, fmt = "d")
plt.title("Confusiion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```





[5.1.4] Wordcloud of top 20 important features from SET 2

In [70]:

```
# Calculate feature importances from decision trees
importances = clf.feature_importances_

# Sort feature importances in descending order
indices = np.argsort(importances)[::-1][:20]

# Rearrange feature names so they match the sorted feature importances
names = tf_idf_vect.get_feature_names()
type(names)
```

Out[70]:

list

In [71]:

```

####https://www.geeksforgeeks.org/generating-word-cloud-python/
# importing all necessary modules
from wordcloud import WordCloud, STOPWORDS
import matplotlib.pyplot as plt
import pandas as pd

comment_words = ' '
stopwords = set(STOPWORDS)

# iterate through the csv file
for val in names:

    # typecaste each val to string
    val = str(val)

    # split the value
    tokens = val.split()

    # Converts each token into lowercase
    for i in range(len(tokens)):
        tokens[i] = tokens[i].lower()

    for words in tokens:
        comment_words = comment_words + words + ' '

wordcloud = WordCloud(width = 800, height = 800, background_color ='white', stopwords = stopwords,
min_font_size = 10).generate(comment_words)

# plot the WordCloud image
plt.figure(figsize = (8, 8), facecolor = None)
plt.imshow(wordcloud)
plt.axis("off")
plt.tight_layout(pad = 0)

plt.show()

```



In [59]:


```
## find hyperparameter using roc_auc score and plot AUC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
base_learners = [40,60,80,100,120]

#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for n in base_learners:
    clf = RandomForestClassifier(n_estimators=n)
    clf.fit(x_train_avg2v, y_train)
    y_train_pred = clf.predict_proba(x_train_avg2v)[:,-1]
    y_test_pred = clf.predict_proba(x_test_avg2v)[:,-1]

    training_scores.append(roc_auc_score(y_train,y_train_pred))
    cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(base_learners, training_scores, label='Train AUC')
plt.plot(base_learners, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("n_estimator: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```



In [59]:

```
## Using GridsearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV

base_learners = [40,60,80,100,120]
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

param_grid = {'n_estimators': base_learners, "max_depth":max_depth_values}
RFC = RandomForestClassifier(max_features='sqrt')
model = GridSearchCV(RFC, param_grid,cv=3 ,pre_dispatch=2,scoring='roc_auc')
model.fit(x_train_avg2v, y_train)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(x_test_avg2v, y_test))
```

Model with best parameters :

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=100, max_features='sqrt', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=120, n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

Accuracy of the model : 0.8997360015989295

In [60]:

```
#### 3D Scatter Plot
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
import numpy as np

x1 = [40,60,80,100,120]
y1 = [1, 5, 10, 50, 100]
z1 = training_scores

x2 = [40,60,80,100,120]
y2 = [1, 5, 10, 50, 100]
z2 = cv_scores
```

In [61]:

```
# https://plot.ly/python/3d-axes/
trace1 = go.Scatter3d(x=x1,y=y1,z=z1, name = 'train')
trace2 = go.Scatter3d(x=x2,y=y2,z=z2, name = 'Cross validation')
data = [trace1, trace2]

layout = go.Layout(scene = dict(
    xaxis = dict(title='n_estimators'),
    yaxis = dict(title='max_depth'),
    zaxis = dict(title='AUC'),))

fig = go.Figure(data=data, layout=layout)
offline.iplot(fig, filename='3d-scatter-colorscale')
```

Observation: Using GridsearchCV max_depth is 100 and n_estimators is 120.

In [60]:

```
# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
from sklearn.metrics import roc_curve, auc
clf = RandomForestClassifier(max_depth=100,n_estimators=120)
clf.fit(x_train_avgw2v, y_train)
y_pred = clf.predict(x_test_avgw2v)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
```


```
# not the predicted outputs

train_fpr, train_tpr, thresholds = roc_curve(y_train, clf.predict_proba(x_train_avgw2v)[: ,1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, clf.predict_proba(x_test_avgw2v)[: ,1])

plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.show()

print("="*100)

from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, clf.predict(x_train_avgw2v)))
print("Test confusion matrix")
print(confusion_matrix(y_test, clf.predict(x_test_avgw2v)))
```



```
=====

Train confusion matrix
[[10993   24]
 [   0 58983]]
Test confusion matrix
[[ 1599  3053]
 [   458 24890]]
```

In [61]:


```
# Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm
```

Out[61]:

```
array([[ 1599,  3053],
       [   458, 24890]])
```

In [62]:

```
# plot confusion matrix to describe the performance of classifier.
import seaborn as sns
class_label = ["negative", "positive"]
df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
sns.heatmap(df_cm, annot = True, fmt = "d")
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```



[5.1.6] Applying Random Forests on TFIDF W2V, SET 4

In [62]:


```
## find hyperparameter using roc_auc score and plot AUC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for d in max_depth_values:
    clf = RandomForestClassifier(max_depth=d)
    clf.fit(x_train_tfidf2v, y_train)
    y_train_pred = clf.predict_proba(x_train_tfidf2v)[:,1]
    y_test_pred = clf.predict_proba(x_test_tfidf2v)[:,1]

    training_scores.append(roc_auc_score(y_train,y_train_pred))
    cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(max_depth_values, training_scores, label='Train AUC')
plt.plot(max_depth_values, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("max_depth: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```



In [63]:

```
## find hyperparameter using roc_auc score and plot AUC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
base_learners = [40,60,80,100,120]
```

```


#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for n in base_learners:
    clf = RandomForestClassifier(n_estimators=n)
    clf.fit(x_train_tfidfw2v, y_train)
    y_train_pred = clf.predict_proba(x_train_tfidfw2v)[: ,1]
    y_test_pred =  clf.predict_proba(x_test_tfidfw2v)[: ,1]

    training_scores.append(roc_auc_score(y_train,y_train_pred))
    cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(base_learners, training_scores, label='Train AUC')
plt.plot(base_learners, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("n_estimator: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```



In [67]:

```

## Using GridsearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV

base_learners = [40,60,80,100,120]
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

param_grid = {'n_estimators': base_learners , "max_depth":max_depth_values}
RFC = RandomForestClassifier(max_features='sqrt')
model = GridSearchCV(RFC, param_grid,cv=3 ,pre_dispatch=2,scoring='roc_auc')
model.fit(x_train_tfidfw2v, y_train)

```

Out[67]:

```

GridSearchCV(cv=3, error_score='raise-deprecating',
  estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
    max_depth=None, max_features='sqrt', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators='warn', n_jobs=None,
    oob_score=False, random_state=None, verbose=0,
    warm_start=False),
  fit_params=None, iid='warn', n_jobs=None,
  param_grid={'n_estimators': [40, 60, 80, 100, 120], 'max_depth': [1, 5, 10, 50, 100, 500, 1
000]},
  pre_dispatch=2, refit=True, return_train_score='warn',
  scoring='roc_auc', verbose=0)

```

In [64]:

```

#### 3D Scatter Plot

```

```

import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
import numpy as np

x1 = [40,60,80,100,120]
y1 = [1, 5, 10, 50, 100]
z1 = training_scores

x2 = [40,60,80,100,120]
y2 = [1, 5, 10, 50, 100]
z2 = cv_scores

```

In [65]:

```

# https://plot.ly/python/3d-axes/
trace1 = go.Scatter3d(x=x1,y=y1,z=z1, name = 'train')
trace2 = go.Scatter3d(x=x2,y=y2,z=z2, name = 'Cross validation')
data = [trace1, trace2]

layout = go.Layout(scene = dict(
    xaxis = dict(title='n_estimators'),
    yaxis = dict(title='max_depth'),
    zaxis = dict(title='AUC'),))

fig = go.Figure(data=data, layout=layout)
offline.iplot(fig, filename='3d-scatter-colorscale')

```

Observation: By observing curve getting high AUC at 100 max_depth and 120 n_estimator.

In [52]:

```

# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
from sklearn.metrics import roc_curve, auc
clf = RandomForestClassifier(max_depth=100,n_estimators=120)
clf.fit(x_train_tfidf2v, y_train)
y_pred = clf.predict(x_test_tfidf2v)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs

train_fpr, train_tpr, thresholds = roc_curve(y_train, clf.predict_proba(x_train_tfidf2v)[:,1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, clf.predict_proba(x_test_tfidf2v)[:,1])

```




```

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.show()

print("="*100)

from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, clf.predict(x_train_tfidfv2v)))
print("Test confusion matrix")
print(confusion_matrix(y_test, clf.predict(x_test_tfidfv2v)))

```



```

=====

Train confusion matrix
[[10782   30]
 [    0 59188]]
Test confusion matrix
[[ 780  4055]
 [ 321 24844]]

```

In [53]:

```

# Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm

```

Out[53]:

```

array([[ 780,  4055],
       [ 321, 24844]])


```


In [54]:

```

# plot confusion matrix to describe the performance of classifier.
import seaborn as sns
class_label = ["negative", "positive"]
df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
sns.heatmap(df_cm, annot = True, fmt = "d")
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```





[5.2] Applying GBDT using XGBOOST

[5.2.1] Applying XGBOOST on BOW, SET 1

In [80]:


```
## find hyperparameter using roc_auc score and plot AUC
from sklearn.metrics import roc_auc_score
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for d in max_depth_values:
    clf = XGBClassifier(max_depth=d)
    clf.fit(x_train_bow, y_train)
    y_train_pred = clf.predict_proba(x_train_bow)[:,1]
    y_test_pred = clf.predict_proba(x_test_bow)[:,1]

    training_scores.append(roc_auc_score(y_train,y_train_pred))
    cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(max_depth_values, training_scores, label='Train AUC')
plt.plot(max_depth_values, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("max_depth: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```



In [81]:

```
## find hyperparameter using roc_auc score and plot AUC
from xgboost import XGBClassifier
from sklearn.metrics import roc_auc_score
base_learners = [40,60,80,100,120]

#empty lists that stores cv scores and training_scores
```

```


cv_scores = []
training_scores = []

#perform k fold cross validation
for n in base_learners:
    clf = XGBClassifier(n_estimators=n)
    clf.fit(x_train_bow, y_train)
    y_train_pred = clf.predict_proba(x_train_bow)[: ,1]
    y_test_pred = clf.predict_proba(x_test_bow)[: ,1]

    training_scores.append(roc_auc_score(y_train,y_train_pred))
    cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(base_learners, training_scores, label='Train AUC')
plt.plot(base_learners, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("n_estimator: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```



In [68]:

```

## Using GridsearchCV
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV

base_learners = [40,60,80,100,120]
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

param_grid = {'n_estimators': base_learners , "max_depth":max_depth_values}
XGB = XGBClassifier(max_features='sqrt')
model = GridSearchCV(XGB, param_grid,cv=3 ,pre_dispatch=2,scoring='roc_auc')
model.fit(x_train_bow, y_train)

```

Out[68]:

```

GridSearchCV(cv=3, error_score='raise-deprecating',
             estimator=XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                                     colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
                                     max_depth=3, max_features='sqrt', min_child_weight=1, missing=None,
                                     n_estimators=100, n_jobs=1, nthread=None,
                                     objective='binary:logistic', random_state=0, reg_alpha=0,
                                     reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
                                     subsample=1),
             fit_params=None, iid='warn', n_jobs=None,
             param_grid={'n_estimators': [40, 60, 80, 100, 120], 'max_depth': [1, 5, 10, 50, 100, 500, 1
000]},
             pre_dispatch=2, refit=True, return_train_score='warn',
             scoring='roc_auc', verbose=0)

```

In [82]:

```

#### 3D Scatter Plot
import plotly.offline as offline

```

```

import plotly.graph_objs as go
offline.init_notebook_mode()
import numpy as np

x1 = [40,60,80,100,120]
y1 = [1, 5, 10, 50, 100]
z1 = training_scores

x2 = [40,60,80,100,120]
y2 = [1, 5, 10, 50, 100]
z2 = cv_scores

```

In [83]:

```

# https://plot.ly/python/3d-axes/
trace1 = go.Scatter3d(x=x1,y=y1,z=z1, name = 'train')
trace2 = go.Scatter3d(x=x2,y=y2,z=z2, name = 'Cross validation')
data = [trace1, trace2]

layout = go.Layout(scene = dict(
    xaxis = dict(title='n_estimators'),
    yaxis = dict(title='max_depth'),
    zaxis = dict(title='AUC'),))

fig = go.Figure(data=data, layout=layout)
offline.iplot(fig, filename='3d-scatter-colorscale')

```

Observation: By GridsearchCV getting high AUC at 3 max_depth and 100 n_estimator.

In [46]:

```

# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
from sklearn.metrics import roc_curve, auc
from xgboost import XGBClassifier
clf = XGBClassifier(max_depth=3,n_estimators=100)
clf.fit(x_train_bow, y_train)
y_pred = clf.predict(x_test_bow)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs

train_fpr, train_tpr, thresholds = roc_curve(y_train, clf.predict_proba(x_train_bow)[:,-1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, clf.predict_proba(x_test_bow)[:,-1])

```

```


test_fpr, test_tpr, thresholds = roc_curve(y_test, clf.predict_proba(x_test_bow)[:,1])

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.show()

print("=="*100)

from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, clf.predict(x_train_bow)))
print("Test confusion matrix")
print(confusion_matrix(y_test, clf.predict(x_test_bow)))

```



```

=====

Train confusion matrix
[[ 2558  8364]
 [   230 58848]]
Test confusion matrix
[[   990  3585]
 [   131 25294]]

```

In [47]:

```

# Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm

```

Out[47]:

```

array([[ 990,  3585],
       [ 131, 25294]])


```


In [48]:

```

# plot confusion matrix to describe the performance of classifier.
import seaborn as sns
class_label = ["negative", "positive"]
df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
sns.heatmap(df_cm, annot = True, fmt = "d")
plt.title("Confusiion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```





[5.2.2] Applying XGBOOST on TFIDF, SET 2

In [84]:


```
## find hyperparameter using roc_auc score and plot AUC
from xgboost import XGBClassifier
from sklearn.metrics import roc_auc_score
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for d in max_depth_values:
    clf = XGBClassifier(max_depth=d)
    clf.fit(x_train_tfidf, y_train)
    y_train_pred = clf.predict_proba(x_train_tfidf)[:,1]
    y_test_pred = clf.predict_proba(x_test_tfidf)[:,1]

    training_scores.append(roc_auc_score(y_train,y_train_pred))
    cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(max_depth_values, training_scores, label='Train AUC')
plt.plot(max_depth_values, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("max_depth: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```



In [85]:


```
## find hyperparameter using roc_auc score and plot AUC
from xgboost import XGBClassifier
from sklearn.metrics import roc_auc_score
base_learners = [40,60,80,100,120]

#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []
```

```
#perform k fold cross validation
for n in base_learners:
    clf = XGBClassifier(n_estimators=n)
    clf.fit(x_train_tfidf, y_train)
    y_train_pred = clf.predict_proba(x_train_tfidf)[:,1]
    y_test_pred = clf.predict_proba(x_test_tfidf)[:,1]

    training_scores.append(roc_auc_score(y_train,y_train_pred))
    cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(base_learners, training_scores, label='Train AUC')
plt.plot(base_learners, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("n_estimator: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```



In [49]:

```
## Using GridsearchCV
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV

base_learners = [40,60,80,100,120]
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

param_grid = {'n_estimators': base_learners , "max_depth":max_depth_values}
XGB = XGBClassifier(max_features='sqrt')
model = GridSearchCV(XGB, param_grid,cv=3 ,pre_dispatch=2,scoring='roc_auc')
model.fit(x_train_tfidf, y_train)
```

Out[49]:

```
GridSearchCV(cv=3, error_score='raise-deprecating',
             estimator=XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                                     colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
                                     max_depth=3, max_features='sqrt', min_child_weight=1, missing=None,
                                     n_estimators=100, n_jobs=1, nthread=None,
                                     objective='binary:logistic', random_state=0, reg_alpha=0,
                                     reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
                                     subsample=1),
             fit_params=None, iid='warn', n_jobs=None,
             param_grid={'n_estimators': [40, 60, 80, 100, 120], 'max_depth': [1, 5, 10, 50, 100, 500, 1
000]},
             pre_dispatch=2, refit=True, return_train_score='warn',
             scoring='roc_auc', verbose=0)
```

Observation: By Observing Curve, getting high AUC at 3 max_depth and 100 n_estimator.

In [86]:

```
#### 3D Scatter Plot
import plotly.offline as offline
import plotly.graph_objs as go
```

```

offline.init_notebook_mode()
import numpy as np

x1 = [40,60,80,100,120]
y1 = [1, 5, 10, 50, 100]
z1 = training_scores

x2 = [40,60,80,100,120]
y2 = [1, 5, 10, 50, 100]
z2 = cv_scores

```

In [87]:

```

# https://plot.ly/python/3d-axes/
trace1 = go.Scatter3d(x=x1,y=y1,z=z1, name = 'train')
trace2 = go.Scatter3d(x=x2,y=y2,z=z2, name = 'Cross validation')
data = [trace1, trace2]

layout = go.Layout(scene = dict(
    xaxis = dict(title='n_estimators'),
    yaxis = dict(title='max_depth'),
    zaxis = dict(title='AUC'),))

fig = go.Figure(data=data, layout=layout)
offline.iplot(fig, filename='3d-scatter-colorscale')

```

In [56]:

```

# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
from sklearn.metrics import roc_curve, auc
clf = XGBClassifier(max_depth=3,n_estimators=100)
clf.fit(x_train_tfidf, y_train)
y_pred = clf.predict(x_test_tfidf)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs

train_fpr, train_tpr, thresholds = roc_curve(y_train, clf.predict_proba(x_train_tfidf)[:,-1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, clf.predict_proba(x_test_tfidf)[:,-1])

plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))


```



```
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.show()

print("="*100)

from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, clf.predict(x_train_tfidf)))
print("Test confusion matrix")
print(confusion_matrix(y_test, clf.predict(x_test_tfidf)))
```



Train confusion matrix

```
[[ 2940  7982]
 [  232 58846]]
```

Test confusion matrix

```
[[ 1145  3430]
 [  129 25296]]
```

In [57]:


```
# Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm
```


Out[57]:

```
array([[ 1145,  3430],
       [  129, 25296]])
```

In [58]:

```
# plot confusion matrix to describe the performance of classifier.
import seaborn as sns
class_label = ["negative", "positive"]
df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
sns.heatmap(df_cm, annot = True, fmt = "d")
plt.title("Confusiion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```





[5.2.3] Applying XGBOOST on AVG W2V, SET 3

In [88]:


```
## find hyperparameter using roc_auc score and plot AUC
from xgboost import XGBClassifier
from sklearn.metrics import roc_auc_score
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for d in max_depth_values:
    clf = XGBClassifier(max_depth=d)
    clf.fit(x_train_avgw2v, y_train)
    y_train_pred = clf.predict_proba(x_train_avgw2v)[:,1]
    y_test_pred = clf.predict_proba(x_test_avgw2v)[:,1]

    training_scores.append(roc_auc_score(y_train,y_train_pred))
    cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(max_depth_values, training_scores, label='Train AUC')
plt.plot(max_depth_values, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("max_depth: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```



In [89]:

```
### find hyperparameter using roc_auc score and plot AUC
from xgboost import XGBClassifier
from sklearn.metrics import roc_auc_score
base_learners = [40,60,80,100,120]

#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for n in base_learners:
    clf = XGBClassifier(n_estimators=n)
```

```

clf.fit(x_train_avgw2v, y_train)
y_train_pred = clf.predict_proba(x_train_avgw2v)[:,-1]
y_test_pred = clf.predict_proba(x_test_avgw2v)[:,-1]


training_scores.append(roc_auc_score(y_train,y_train_pred))
cv_scores.append(roc_auc_score(y_test, y_test_pred))

```

```

#plot cross-validated score, training score vs alpha
plt.plot(base_learners, training_scores, label='Train AUC')
plt.plot(base_learners, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("n_estimator: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```



In [69]:

```

## Using GridsearchCV
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV

base_learners = [40,60,80,100,120]
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

param_grid = {'n_estimators': base_learners , "max_depth":max_depth_values}
RFC = XGBClassifier(max_features='sqrt')
model = GridSearchCV(XGB, param_grid,cv=3 ,pre_dispatch=2,scoring='roc_auc')
model.fit(x_train_avgw2v, y_train)

```

Out[69]:

```

GridSearchCV(cv=3, error_score='raise-deprecating',
  estimator=XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
    colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
    max_depth=3, max_features='sqrt', min_child_weight=1, missing=None,
    n_estimators=100, n_jobs=1, nthread=None,
    objective='binary:logistic', random_state=0, reg_alpha=0,
    reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
    subsample=1),
  fit_params=None, iid='warn', n_jobs=None,
  param_grid={'n_estimators': [40, 60, 80, 100, 120], 'max_depth': [1, 5, 10, 50, 100, 500, 1
000]},
  pre_dispatch=2, refit=True, return_train_score='warn',
  scoring='roc_auc', verbose=0)

```

In [90]:

```

#### 3D Scatter Plot
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
import numpy as np

x1 = [40,60,80,100,120]
v1 = [1, 5, 10, 50, 100]

```

```

x1 = [1, 5, 10, 50, 100]
z1 = training_scores

x2 = [40, 60, 80, 100, 120]
y2 = [1, 5, 10, 50, 100]
z2 = cv_scores

```

In [91]:

```

# https://plot.ly/python/3d-axes/
trace1 = go.Scatter3d(x=x1,y=y1,z=z1, name = 'train')
trace2 = go.Scatter3d(x=x2,y=y2,z=z2, name = 'Cross validation')
data = [trace1, trace2]

layout = go.Layout(scene = dict(
    xaxis = dict(title='n_estimators'),
    yaxis = dict(title='max_depth'),
    zaxis = dict(title='AUC'),))

fig = go.Figure(data=data, layout=layout)
offline.iplot(fig, filename='3d-scatter-colorscale')

```

Observation: By observing curve getting high AUC at 3 max_depth and 100 n_estimator.

In [53]:

```

# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
from sklearn.metrics import roc_curve, auc
clf = XGBClassifier(max_depth=3,n_estimators=100)
clf.fit(x_train_avg2v, y_train)
y_pred = clf.predict(x_test_avg2v)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs

train_fpr, train_tpr, thresholds = roc_curve(y_train, clf.predict_proba(x_train_avg2v)[:,1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, clf.predict_proba(x_test_avg2v)[:,1])


plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")

```

```
plt.title("ERROR PLOTS")
plt.show()

print("="*100)

from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, clf.predict(x_train_avgw2v)))
print("Test confusion matrix")
print(confusion_matrix(y_test, clf.predict(x_test_avgw2v)))
```



```
=====

Train confusion matrix
[[ 4798  6124]
 [ 1331 57747]]
Test confusion matrix
[[ 1844  2731]
 [  663 24762]]
```

In [54]:


```
# Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm
```


Out[54]:

```
array([[ 1844,  2731],
       [   663, 24762]])
```

In [55]:

```
# plot confusion matrix to describe the performance of classifier.
import seaborn as sns
class_label = ["negative", "positive"]
df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
sns.heatmap(df_cm, annot = True, fmt = "d")
plt.title("Confusiion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```





[5.2.4] Applying XGBOOST on TFIDF W2V, SET 4

In [92]:


```
## find hyperparameter using roc_auc score and plot AUC
from xgboost import XGBClassifier
from sklearn.metrics import roc_auc_score
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for d in max_depth_values:
    clf = XGBClassifier(max_depth=d)
    clf.fit(x_train_tfidf2v, y_train)
    y_train_pred = clf.predict_proba(x_train_tfidf2v)[:,-1]
    y_test_pred = clf.predict_proba(x_test_tfidf2v)[:,-1]

    training_scores.append(roc_auc_score(y_train,y_train_pred))
    cv_scores.append(roc_auc_score(y_test, y_test_pred))

#plot cross-validated score, training score vs alpha
plt.plot(max_depth_values, training_scores, label='Train AUC')
plt.plot(max_depth_values, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("max_depth: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```



In [93]:


```
## find hyperparameter using roc_auc score and plot AUC
from xgboost import XGBClassifier
from sklearn.metrics import roc_auc_score
base_learners = [40,60,80,100,120]

#empty lists that stores cv scores and training_scores
cv_scores = []
training_scores = []

#perform k fold cross validation
for n in base_learners:
    clf = XGBClassifier(n_estimators=n)
    clf.fit(x_train_tfidf2v, y_train)
    y_train_pred = clf.predict_proba(x_train_tfidf2v)[:,-1]
    y_test_pred = clf.predict_proba(x_test_tfidf2v)[:,-1]
```

```
training_scores.append(roc_auc_score(y_train,y_train_pred))
cv_scores.append(roc_auc_score(y_test, y_test_pred))
```

```
#plot cross-validated score, training score vs alpha
plt.plot(base_learners, training_scores, label='Train AUC')
plt.plot(base_learners, cv_scores, label='CV AUC')
plt.legend()
plt.xlabel("n_estimator: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```



In []:

```
## Using GridsearchCV
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV

base_learners = [40,60,80,100,120]
max_depth_values = [1, 5, 10, 50, 100, 500, 1000]

param_grid = {'n_estimators': base_learners , "max_depth":max_depth_values}
RFC = XGBClassifier(max_features='sqrt')
model = GridSearchCV(XGB, param_grid,cv=3 ,pre_dispatch=2,scoring='roc_auc')
model.fit(x_train_tfdfw2v, y_train)
```

Out[]:

```
GridSearchCV(cv=3, error_score='raise-deprecating',
  estimator=XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
  colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
  max_depth=3, max_features='sqrt', min_child_weight=1, missing=None,
  n_estimators=100, n_jobs=1, nthread=None,
  objective='binary:logistic', random_state=0, reg_alpha=0,
  reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
  subsample=1),
  fit_params=None, iid='warn', n_jobs=None,
  param_grid={'n_estimators': [40, 60, 80, 100, 120], 'max_depth': [1, 5, 10, 50, 100, 500, 1
000]},
  pre_dispatch=2, refit=True, return_train_score='warn',
  scoring='roc_auc', verbose=0)
```

In [94]:

```
#### 3D Scatter Plot
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
import numpy as np

x1 = [40,60,80,100,120]
y1 = [1, 5, 10, 50, 100]
z1 = training_scores

x2 = [40,60,80,100,120]
```

```
y2 = [1, 5, 10, 50, 100]
z2 = cv_scores
```

In [95]:

```
# https://plot.ly/python/3d-axes/
trace1 = go.Scatter3d(x=x1,y=y1,z=z1, name = 'train')
trace2 = go.Scatter3d(x=x2,y=y2,z=z2, name = 'Cross validation')
data = [trace1, trace2]

layout = go.Layout(scene = dict(
    xaxis = dict(title='n_estimators'),
    yaxis = dict(title='max_depth'),
    zaxis = dict(title='AUC'),))

fig = go.Figure(data=data, layout=layout)
offline.iplot(fig, filename='3d-scatter-colorscale')
```

Observation: By observing curve getting high AUC at 3 max_depth and 100 n_estimator.

In [50]:

```
# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
from sklearn.metrics import roc_curve, auc
clf = XGBClassifier(max_depth=3,n_estimators=100)
clf.fit(x_train_tfidf2v, y_train)
y_pred = clf.predict(x_test_tfidf2v)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs

train_fpr, train_tpr, thresholds = roc_curve(y_train, clf.predict_proba(x_train_tfidf2v)[:,1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, clf.predict_proba(x_test_tfidf2v)[:,1])

plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.show()


print(" "*100)
```



```

from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, clf.predict(x_train_tfidfw2v)))
print("Test confusion matrix")
print(confusion_matrix(y_test, clf.predict(x_test_tfidfw2v)))

```



Train confusion matrix

```
[[ 2579  8343]
 [  918 58160]]
```

Test confusion matrix

```
[[  987  3588]
 [  429 24996]]
```

In [51]:

```

# Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm

```

Out[51]:


```
array([[ 987,  3588],
       [  429, 24996]])
```

In [52]:

```

# plot confusion matrix to describe the performance of classifier.
import seaborn as sns
class_label = ["negative", "positive"]
df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
sns.heatmap(df_cm, annot = True, fmt = "d")
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```



[6] Conclusions

In []:

```
# Please compare all your models using Prettytable library
```

In [55]:

```
from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Vectorizer", "Model", "max_depth", "n_estimator", "AUC"]

x.add_row(["BOW", "RF", "1000", "120", "92.83%"])
x.add_row(["TFIDF", "RF", "500", "40", "92.93%"])
x.add_row(["AvgW2V", "RF", "100", "120", "89.89%"])
x.add_row(["TFIDF W2V", "RF", "100", "120", "84.54%"])
x.add_row(["BOW", "XGB", "3", "100", "87.63%"])
x.add_row(["TFIDF", "XGB", "3", "100", "88.43%"])
x.add_row(["AvgW2V", "XGB", "3", "100", "89.89%"])
x.add_row(["TFIDF W2V", "XGB", "3", "100", "84.30%"])

print(x)
```

Vectorizer	Model	max_depth	n_estimator	AUC
BOW	RF	1000	120	92.83%
TFIDF	RF	500	40	92.93%
AvgW2V	RF	100	120	89.89%
TFIDF W2V	RF	100	120	84.54%
BOW	XGB	3	100	87.63%
TFIDF	XGB	3	100	88.43%
AvgW2V	XGB	3	100	89.89%
TFIDF W2V	XGB	3	100	84.30%

1) RandomForest with TFIDF getting high accuracy. 2) GridsearchCV takes a long time to train. 3) In XGBoost classifier getting same hyperparameters for all 4 featurization techniques. 4) In both RandomForest and XGBoost classifier, with AVGW2v and TFIDFW2V getting same accuracy. 5) In RandomForest Classifier with TFIDFW2V, Hyperparameters are not getting with GridsearchCV.